

STAT 206 Lab 2

Due Monday, October 16, 5:00 PM

General instructions for labs: You are encouraged to work in pairs to complete the lab. Labs must be completed as an R Markdown file. Be sure to include your lab partner (if you have one) and your own name in the file. Give the commands to answer each question in its own code block, which will also produce plots that will be automatically embedded in the output file. Each answer must be supported by written statements as well as any code used.

Agenda: Manipulating data frames; practicing iteration; practicing re-writing code; checking how reliable random methods are.

Part I – Data Frames

R includes a number of pre-specified data objects as part of its default installation. We will load and manipulate one of these, a data frame of 93 cars with model year 1993. Begin by ensuring that you can load this data with the commands

```
library(MASS)
data(Cars93)
```

Begin by examining the data frame with the command `View(Cars93)` to understand the underlying object. You will need to use functions and other commands to extract elements for this assignment.

```
# Class displays what type a given argument is
```

```
class(Cars93)
```

```
## [1] "data.frame"
```

```
#View puts the data frame into its own table with header
```

```
View(Cars93)
```

1. Obtain a `summary()` of the full data structure. Can you tell from this how many rows are in the data? If so, say how; if not, use another method to obtain the number of rows.

```
#the summary function returns some descriptive statistics on from the data frame. It includes the "5 num  
#the 5 number summary includes the min value, q1, median, q3 and max value  
#From this data we can get an idea of how many rows and columns. We can count the number of attributes  
#summary returns and we can know the amount of columns, and for the rows we can find an attribute that  
#an index to find the number of rows. in this case we see both Model and Make display 6 values and shows  
#My guess would be that there is 93 rows  
#It is probably much better to use a couple functions to accomplish this and display the number of colum
```

```
summary(Cars93)
```

##	Manufacturer	Model	Type	Min.Price	Price
##	Chevrolet: 8	100 : 1	Compact:16	Min. : 6.70	Min. : 7.40
##	Ford : 8	190E : 1	Large :11	1st Qu.:10.80	1st Qu.:12.20
##	Dodge : 6	240 : 1	Midsize:22	Median :14.70	Median :17.70
##	Mazda : 5	300E : 1	Small :21	Mean :17.13	Mean :19.51
##	Pontiac : 5	323 : 1	Sporty :14	3rd Qu.:20.30	3rd Qu.:23.30
##	Buick : 4	535i : 1	Van : 9	Max. :45.40	Max. :61.90

```

## (Other) :57 (Other):87
## Max.Price MPG.city MPG.highway AirBags
## Min. : 7.9 Min. :15.00 Min. :20.00 Driver & Passenger:16
## 1st Qu.:14.7 1st Qu.:18.00 1st Qu.:26.00 Driver only :43
## Median :19.6 Median :21.00 Median :28.00 None :34
## Mean :21.9 Mean :22.37 Mean :29.09
## 3rd Qu.:25.3 3rd Qu.:25.00 3rd Qu.:31.00
## Max. :80.0 Max. :46.00 Max. :50.00
##
## DriveTrain Cylinders EngineSize Horsepower RPM
## 4WD :10 3 : 3 Min. :1.000 Min. : 55.0 Min. :3800
## Front:67 4 :49 1st Qu.:1.800 1st Qu.:103.0 1st Qu.:4800
## Rear :16 5 : 2 Median :2.400 Median :140.0 Median :5200
## 6 :31 Mean :2.668 Mean :143.8 Mean :5281
## 8 : 7 3rd Qu.:3.300 3rd Qu.:170.0 3rd Qu.:5750
## rotary: 1 Max. :5.700 Max. :300.0 Max. :6500
##
## Rev.per.mile Man.trans.avail Fuel.tank.capacity Passengers
## Min. :1320 No :32 Min. : 9.20 Min. :2.000
## 1st Qu.:1985 Yes:61 1st Qu.:14.50 1st Qu.:4.000
## Median :2340 Median :16.40 Median :5.000
## Mean :2332 Mean :16.66 Mean :5.086
## 3rd Qu.:2565 3rd Qu.:18.80 3rd Qu.:6.000
## Max. :3755 Max. :27.00 Max. :8.000
##
## Length Wheelbase Width Turn.circle
## Min. :141.0 Min. : 90.0 Min. :60.00 Min. :32.00
## 1st Qu.:174.0 1st Qu.: 98.0 1st Qu.:67.00 1st Qu.:37.00
## Median :183.0 Median :103.0 Median :69.00 Median :39.00
## Mean :183.2 Mean :103.9 Mean :69.38 Mean :38.96
## 3rd Qu.:192.0 3rd Qu.:110.0 3rd Qu.:72.00 3rd Qu.:41.00
## Max. :219.0 Max. :119.0 Max. :78.00 Max. :45.00
##
## Rear.seat.room Luggage.room Weight Origin
## Min. :19.00 Min. : 6.00 Min. :1695 USA :48
## 1st Qu.:26.00 1st Qu.:12.00 1st Qu.:2620 non-USA:45
## Median :27.50 Median :14.00 Median :3040
## Mean :27.83 Mean :13.89 Mean :3073
## 3rd Qu.:30.00 3rd Qu.:15.00 3rd Qu.:3525
## Max. :36.00 Max. :22.00 Max. :4105
## NA's :2 NA's :11
##
## Make
## Acura Integra: 1
## Acura Legend : 1
## Audi 100 : 1
## Audi 90 : 1
## BMW 535i : 1
## Buick Century: 1
## (Other) :87

```

```
nrow(Cars93)
```

```
## [1] 93
```

```
ncol(Cars93)
```

```
## [1] 27
```

nrow and ncol will display the number of rows and columns in a dataframe, respectively. This is a mu

2. What is the mean price of a car with a rear-wheel drive train?

*#First step is to get a slice of this data frame to extract only rows that have 'Rear' for attribute Dr
#We also want to get the price values associated with the above values
#we can use the subset function for an intuitive way to obtain this slice*

```
rwd_avg_p = subset(Cars93, DriveTrain == 'Rear', select = c('Price', 'DriveTrain'))
```

#Finally, we use the mean function but only on the price values in our subset

```
mean(rwd_avg_p[, 'Price'])
```

```
## [1] 28.95
```

3. What is the minimum horsepower of all cars with capacity for 7 passengers? With a capacity of at least 6 passengers?

*#Again the subset function will be used to grab a slice of the data frame that consists of horse power
#first will be a slice for when passengers = 7*

```
cap_hp = subset(Cars93, Passengers == 7, select = c('Passengers', 'Horsepower'))
```

#the min function is called on the attribute Horsepower for the subset data

```
min(cap_hp[, 'Horsepower'])
```

```
## [1] 109
```

#Next we will get a slice for when there are atleast 6 passengers

```
cap_hp2 = subset(Cars93, Passengers >= 6, select = c('Passengers', 'Horsepower'))
```

#now we will use the min function again to get the new minimum of this slice

```
min(cap_hp2[, 'Horsepower'])
```

```
## [1] 100
```

4. Assuming that these cars are exactly as fuel efficient as this table indicates, find the cars that have the maximum, minimum and median distance travelable for highway driving. You will need at least two columns to work this out; why those two?

*# To obtain which cars have the best, worst and middle hwy mpg we need to first know what these values
#we can do this using the max, min, and median functions on the Cars93 data, on the hwy mpg column
note this info is available in the summary above*

```
maxmpg = max(Cars93[, 'MPG.highway'])  
maxmpg
```

```
## [1] 50
```

```
medianmpg = median(Cars93[, 'MPG.highway'])  
medianmpg
```

```
## [1] 28
minmpg = min(Cars93[, 'MPG.highway'])
minmpg

## [1] 20
# Now we can create slices of the data frame for Cars that have the above values

maxhighway = subset(Cars93, MPG.highway == maxmpg, select = c('MPG.highway', 'Make'))
maxhighway

##      MPG.highway      Make
## 39             50 Geo Metro

medianhighway = subset(Cars93, MPG.highway == medianmpg, select = c('Make', 'MPG.highway'))
medianhighway

##              Make MPG.highway
## 7      Buick LeSabre      28
## 14    Chevrolet Camaro      28
## 20    Chrylser Concorde      28
## 21    Chrysler LeBaron      28
## 30      Eagle Vision      28
## 71 Oldsmobile Eighty-Eight      28
## 75    Pontiac Firebird      28
## 77    Pontiac Bonneville      28
## 92          Volvo 240      28
## 93          Volvo 850      28

minhighway = subset(Cars93, MPG.highway == minmpg, select = c('Make', 'MPG.highway'))
minhighway

##              Make MPG.highway
## 17 Chevrolet Astro      20
## 36   Ford Aerostar      20
```

Part II – Reproducibility and Functions

Some of the lectures have included examples of planning production for a factory that turns steel and labor into cars and trucks. Below is a piece of code that optimizes the factory's output (roughly) given the available resources, using a `repeat` loop. It's embedded in a function to make it easier for you to run.

```
factory.function <- function (cars.output=1, trucks.output=1) {
  factory <- matrix(c(40,1,60,3),nrow=2,
    dimnames=list(c("labor", "steel"), c("cars", "trucks")))
  available <- c(1600,70); names(available) <- rownames(factory)
  slack <- c(8,1); names(slack) <- rownames(factory)
  output <- c(cars.output, trucks.output); names(output) <- colnames(factory)

  passes <- 0 # How many times have we been around the loop?
  repeat {
    passes <- passes + 1
    needed <- factory %*% output # What do we need for that output level?
    # If we're not using too much, and are within the slack, we're done
    if (all(needed <= available) &&
```

```

    all((available - needed) <= slack)) {
      break()
    }
    # If we're using too much of everything, cut back by 10%
    if (all(needed > available)) {
      output <- output * 0.9
      next()
    }
    # If we're using too little of everything, increase by 10%
    if (all(needed < available)) {
      output <- output * 1.1
      next()
    }
    # If we're using too much of some resources but not others, randomly
    # tweak the plan by up to 10%
    # runif == Random number, UNIFormly distributed, not "run if"
    output <- output * (1+runif(length(output),min=-0.1,max=0.1))
  }

  return(output)
}

```

5. Run the function above with the command

```
factory.function()
```

```
##      cars      trucks
## 10.28648 19.69488
```

to obtain a default output value, starting from a very low initial planned output. What is the final output capacity obtained?

6. Repeat this four more times to obtain new output values. Do these answers differ from each other? If so why? If not, why not?

```
factory.function()
```

```
##      cars      trucks
## 9.949649 19.959082
```

```
factory.function()
```

```
##      cars      trucks
## 10.08165 19.91860
```

```
factory.function()
```

```
##      cars      trucks
## 10.47071 19.64083
```

```
factory.function()
```

```
##      cars      trucks
## 10.25063 19.75981
```

#All of the answers are relatively close to one another however, all of them are different. This is due to the function itself. It even says that it will "randomly tweak the plan by up to 10%" This means that each time the function will return a different output

7. Right now, the number of `passes` is a value held within the function itself and not shared. Change the code so that the number of `passes` will be returned at the end of the function, as well as the final output.

#The return(output) at the end was commented out so I could use a print statement to display the output

```
factory.function <- function (cars.output=1, trucks.output=1) {
  factory <- matrix(c(40,1,60,3),nrow=2,
    dimnames=list(c("labor","steel"),c("cars","trucks")))
  available <- c(1600,70); names(available) <- rownames(factory)
  slack <- c(8,1); names(slack) <- rownames(factory)
  output <- c(cars.output, trucks.output); names(output) <- colnames(factory)

  passes <- 0 # How many times have we been around the loop?
  repeat {
    passes <- passes + 1
    needed <- factory %*% output # What do we need for that output level?
    # If we're not using too much, and are within the slack, we're done
    if (all(needed <= available) &&
      all((available - needed) <= slack)) {
      break()
    }
    # If we're using too much of everything, cut back by 10%
    if (all(needed > available)) {
      output <- output * 0.9
      next()
    }
    # If we're using too little of everything, increase by 10%
    if (all(needed < available)) {
      output <- output * 1.1
      next()
    }
    # If we're using too much of some resources but not others, randomly
    # tweak the plan by up to 10%
    # runif == Random number, UNIFormly distributed, not "run if"
    output <- output * (1+runif(length(output),min=-0.1,max=0.1))

  }

  #return(output)
  print(output) + print(passes)
}
```

factory.function()

```
##      cars   trucks
## 10.54919 19.61909
## [1] 2239
```

8. Now, set the initial output levels to 30 cars and 20 trucks and run the code. What is the final output plan (output)? What is the final demand for resources (needed)? Is the plan within budget and within the slack? How many iterations did it take to converge (passes)? For all but output you will need to either print this message out deliberately, or return an object that contains all the quantities you want.

the code was copied from above but in the print statement we are printing out 'needed' to see how much material was consumed for the output. In this case the output was almost 10 cars and 20 trucks

```
factory.function <- function (cars.output=1, trucks.output=1) {
  factory <- matrix(c(40,1,60,3),nrow=2,
    dimnames=list(c("labor","steel"),c("cars","trucks")))
  available <- c(1600,70); names(available) <- rownames(factory)
  slack <- c(8,1); names(slack) <- rownames(factory)
  output <- c(cars.output, trucks.output); names(output) <- colnames(factory)

  passes <- 0 # How many times have we been around the loop?
  repeat {
    passes <- passes + 1
    needed <- factory %*% output # What do we need for that output level?
    # If we're not using too much, and are within the slack, we're done
    if (all(needed <= available) &&
      all((available - needed) <= slack)) {
      break()
    }
    # If we're using too much of everything, cut back by 10%
    if (all(needed > available)) {
      output <- output * 0.9
      next()
    }
    # If we're using too little of everything, increase by 10%
    if (all(needed < available)) {
      output <- output * 1.1
      next()
    }
    # If we're using too much of some resources but not others, randomly
    # tweak the plan by up to 10%
    # runif == Random number, UNIFormly distributed, not "run if"
    output <- output * (1+runif(length(output),min=-0.1,max=0.1))

  }

  #return(output)
  print(output) + print(passes) + print(needed)
}
```

```
factory.function(cars.output = 1, trucks.output = 20)
```

```
##      cars      trucks
## 10.26458 19.79964
## [1] 1146
##           [,1]
## labor 1598.56186
## steel  69.66351
```